

Intermediate SQL

Joseph Hallett

February 13, 2023



Last time...

We introduced SQL as a language for querying databases

- ▶ How to create tables
- ▶ How to add and delete data
- ▶ How to run basic queries

This time...

More advanced SQL!

- ▶ More features, more function
- ▶ Other joins
- ▶ NULL

NULL is nothing

There is a special value in SQL to represent missing data: NULL.

- ▶ But they're *pretty much always* a bad idea
- ▶ The logic for comparing them is pretty whacky

NULL = NULL?

Lets say we have a database with the following table:

Person	Fruit
Joseph	Lime
Matt	Apple
Partha	

Lets find everyone who we know what their favourite fruit is!

```
SELECT * FROM fruit WHERE fruit <> NULL;
```

Err..., lets try the opposite?

```
SELECT * FROM fruit WHERE fruit = NULL;
```

Err what?

```
SELECT * FROM fruit WHERE fruit LIKE '%';
```

Person	Fruit
Joseph	Lime
Matt	Apple

So...

```
SELECT * FROM fruit WHERE fruit NOT LIKE '%';
```

NULL is weird...

Because **NULL means attribute missing...**

- ▶ The results of comparing with it are ~~just plain stupid~~ somewhat unexpected

The simple solution is to declare *everything* as NOT NULL

- ▶ And use a higher normal form (5NF) then you'll find they almost entirely disappear

Otherwise you have to memorise a bunch of ~~stupid~~ special comparators

```
SELECT * FROM fruit WHERE fruit IS NULL;
```

Person	Fruit
Partha	

```
SELECT * FROM fruit WHERE fruit IS NOT NULL;
```

Person	Fruit
Joseph	Lime
Matt	Apple

Tricky joins

Clearly testing for equality when NULL is problematic.

- So what happens when you want to join two tables together with NULL's in them

Person	Fruit
Joseph	Lime
Matt	Apple
Partha	

Fruit	Dish
Apple	Apple crumble
Banana	Banana split
Cherry	
Lime	Daiquiri

What's my favourite food?

So what might make a nice dish for each of your lecturers?

- ▶ (A **NATURAL JOIN** is like a regular JOIN but assumes same named columns ought to be equal).

Person	Fruit	Dish
Joseph	Lime	Daiquiri
Matt	Apple	Apple crumble

But what about poor *Partha*? How do we get him to appear in our table?

LEFT and RIGHT JOIN

When doing our previous JOIN we wanted only rows that matched...

- ▶ Technically called an **INNER JOIN**...

Sometimes we're okay with the database sticking NULL in if we want to keep columns where a join *can't* be made...

```
SELECT person, fruit.fruit, dish  
FROM fruit  
LEFT JOIN recipes  
ON fruit.fruit = recipes.fruit;
```

Person	Fruit	Dish
Joseph	Lime	Daiquiri
Matt	Apple	Apple crumble
Partha		

RIGHT JOIN

A RIGHT JOIN is like a left join but the other way round...

```
SELECT fruit.fruit, dish, person
FROM fruit
RIGHT JOIN recipes
ON fruit.fruit = recipes.fruit;
```

Fruit	Dish	Person
Lime	Daiquiri	Joseph
Apple	Apple crumble	Matt
	Banana split	

Where has the Banana gone?!

```
SELECT recipes.fruit, dish, person
FROM fruit
RIGHT JOIN recipes
ON fruit.fruit = recipes.fruit;
```

Fruit	Dish	Person
Lime	Daiquiri	Joseph
Apple	Apple crumble	Matt
Banana	Banana split	
Cherry		

(Or just NATURAL JOIN and it'll *usually* take care of it...)

```
SELECT fruit, dish, person
FROM fruit
RIGHT NATURAL JOIN recipes;
```

Fruit	Dish	Person
Lime	Daiquiri	Joseph
Apple	Apple crumble	Matt
Banana	Banana split	
Cherry		

One more JOIN!

What if we want to do a LEFT and a RIGHT JOIN at the same time?

```
SELECT *  
FROM fruit  
FULL OUTER NATURAL JOIN recipes;
```

Person	Fruit	Dish
Joseph	Lime	Daiquiri
Matt	Apple	Apple crumble
Partha	Banana	Banana split
	Cherry	

What about statistic functions?

In the last lecture we introduced `COUNT` as a way of counting how many things exist?

- How many different fruits are in the outer joined table?

```
SELECT *  
FROM fruit  
FULL OUTER NATURAL JOIN recipes;
```

Person	Fruit	Dish
Joseph	Lime	Daiquiri
Matt	Apple	Apple crumble
Partha	Banana	Banana split
	Cherry	

```
SELECT COUNT(fruit)  
FROM fruit  
FULL OUTER NATURAL JOIN recipes
```

COUNT(fruit)
4

...So it looks like COUNT ignores NULL

Other statistics...

Lets rank fruits!

Fruit	Stars
Apple	0
Banana	4
Cherry	
Lime	5

```
SELECT AVG(stars) AS Average FROM ranking;
```

Average
3.0

```
SELECT SUM(stars)/COUNT(fruit) AS Average  
FROM ranking;
```

Average
2

Remember computers are *awful*

- ▶ Multiply count by 1.0 to "fix"?
- ▶ Also number of stars is ordinal data so the *mean* shouldn't be used anyway...

将数据按照一定的顺序排列

What about standard deviation?

The standard deviation is how far something deviates *on average* from the *mean*.

```
SELECT SQRT(AVG(Deviation)) AS STDDEV
FROM (
  SELECT Fruit, Stars, Mean,
         (Stars-Mean)*(Stars-Mean) AS Deviation
  FROM ranking JOIN (
    SELECT AVG(stars) AS Mean
    FROM ranking
  )
  WHERE stars IS NOT NULL
);
```

STDDEV
2.16024689946929

You can nest queries inside one another (subqueries!)

- ▶ This is a recipe for making your SQL *slow*
- ▶ Maybe just use SQL for data retrieval and leave complex stats to statistical programming languages?

So that's SQL!

Tips for using it?

- ▶ Don't overcomplicate things!
- ▶ Normal forms make things simpler!
- ▶ Avoid NULL like the plague